

CSE331 Introduction to Algorithms

Lecture 22

Optimal Binary Search Trees

Antoine Vigneron
antoine@unist.ac.kr

Ulsan National Institute of Science and Technology

July 23, 2021

- 1 Introduction
- 2 Binary search trees
- 3 Optimal binary search trees
- 4 Structure of an optimal BST
- 5 Recurrence relation
- 6 Computing the optimal cost
- 7 Computing an optimal BST
- 8 Conclusion


Introduction

- In this lecture, we will show how to compute optimal *binary search trees* (BST) by dynamic programming.
- **Reference:** Section 12 (p. 286) and Section 15.5 of the textbook (p. 397)

[Introduction to Algorithms](#) by Cormen, Leiserson, Rivest and Stein.

Subtrees of a BST


(from previous lecture)




- BST with set of keys {5, 9, 10, 15, 16, 18, 20, 21, 22, 25, 28, 30}.

BST Insertion Orders (from previous lecture)


- The shape of a BST depends on the order of insertions.




$1 \rightarrow 2 \rightarrow 3$




$1 \rightarrow 3 \rightarrow 2$




$2 \rightarrow 1 \rightarrow 3$



$2 \rightarrow 3 \rightarrow 1$



$3 \rightarrow 1 \rightarrow 2$



$3 \rightarrow 2 \rightarrow 1$

Optimal Binary Search Trees

- We are given a set of keys $k_1 < k_2 < \dots < k_n$.
- We want to build a BST on this set of keys, so as to answer search queries efficiently.
- We know in advance the probability distribution of the queries.


Example

i	1	2	3	4	5
p_i	0.1	0.2	0.1	0.2	0.4

Problem

Given p_1, \dots, p_n , build a BST with keys k_1, \dots, k_n that minimizes the average search time.


Example



- Probabilities:

i	1	2	3	4	5
p_i	0.1	0.2	0.1	0.2	0.4
- Which tree is better?
- Intuitively, T is better because k_5 is the most frequent key we search, and the search path to k_5 is shorter.

Depth of a Node



Definition

The *depth* of a node v in tree T is the number of edges in the path from the root to v . We denote it by $\text{depth}_T(v)$. As we assume that the keys are distinct, we write $\text{depth}_T(k_i)$ for the depth of the node with key k_i .

Problem Formulation

- We assume that the search time for a node v is proportional to $\text{depth}_T(v) + 1$, which is the number of nodes on the path from the root to v .
- Then the expected search time in a tree T is proportional to

$$\sum_{i=1}^n p_i(\text{depth}_T(k_i) + 1) = \sum_{i=1}^n p_i + \sum_{i=1}^n p_i \text{depth}_T(k_i)$$

- So we want to find a tree T that minimizes

$$\text{cost}(T) = \sum_{i=1}^n p_i + \sum_{i=1}^n p_i \text{depth}_T(k_i).$$

- Remark: We will not assume that $\sum_{i=1}^n p_i = 1$, even though in the original problem it is true because these numbers are probabilities.

Problem Formulation

Problem

Given keys $k_1 < k_2 < \dots < k_n$ and weights p_1, \dots, p_n , the **optimal BST problem** is to find a binary search tree T^* over k_1, \dots, k_n such that

$$\text{cost}(T^*) = \sum_{i=1}^n p_i + \sum_{i=1}^n p_i \text{depth}_{T^*}(k_i)$$


is minimized. The optimal cost is denoted

$$\text{cost}^*(p_1, \dots, p_n) = \text{cost}(T^*).$$

Brute Force Approach

- A brute force approach would enumerate all possible BSTs over k_1, \dots, k_n , compute their costs, and return the smallest.
- It would take exponential time because there is an exponential number of such BSTs.
 - ▶ Proof?
- So we will try dynamic programming.

Structure of an Optimal BST



Lemma

If T^* is an optimal BST, then the left and right subtrees at his root k_r are optimal BSTs over weights (p_1, \dots, p_{r-1}) and (p_{r+1}, \dots, p_n) , respectively.

- Proof: see exercise set.

Recurrence Relation

- Minimizing over all possible choices for the root, we obtain the following recurrence relation.

Corollary

Let $c[i, j] = \text{cost}^*(p_i, \dots, p_j)$. Then

$$c[i, j] = \left(\sum_{k=i}^j p_k \right) + \min_{i \leq r \leq j} \left(c[i, r-1] + c[r+1, j] \right),$$

where $c[i, i-1] = c[j+1, j] = 0$ is the cost of an empty subtree.

Recursive Solution

Pseudocode

```
1: procedure COST*( $p, i, j$ )
2:   if  $j = i - 1$  then                                ▷ base case
3:     return 0
4:   result  $\leftarrow \infty$ 
5:   for  $r \leftarrow i, j$  do                          ▷ minimization
6:     result  $\leftarrow \min \left( \text{result}, \text{COST}^*(p, i, r - 1) + \text{COST}^*(p, r + 1, j) \right)$ 
7:   for  $k \leftarrow i, j$  do
8:     result  $\leftarrow \text{result} + p[k]$                   ▷ adding  $\sum_{k=i}^j p_k$ 
9:   return result
```

- What is the problem with this algorithm?
- It runs in exponential time (see exercise set).
- So we will use dynamic programming


Dynamic Programming: First Attempt

Pseudocode

```
1: procedure DPCost( $p$ ,  $1, n$ )
2:    $c[1 \dots n + 1, 0 \dots n] \leftarrow$  new array
3:   for  $i \leftarrow 1, n + 1$  do                                ▷ base cases
4:      $c[i, i - 1] \leftarrow 0$ 
5:   for  $i \leftarrow 1, n$  do                                ▷ recursive case
6:     for  $j \leftarrow i, n$  do
7:        $c[i, j] \leftarrow \infty$ 
8:       for  $r \leftarrow i, j$  do
9:          $c[i, j] \leftarrow \min(c[i, j], c[i, r - 1] + c[r + 1, j])$ 
10:      for  $k \leftarrow i, j$  do
11:         $c[i, j] \leftarrow c[i, j] + p[k]$ 
12:   return  $c[1, n]$ 
```

Dynamic Programming: First Attempt

- What is the problem with this algorithm?
- It does not compute the values $c[i, j]$ in the right order.
- For instance, $c[1, 2]$ is obtained from $c[1, 1]$ and $c[2, 2]$, but $c[2, 2]$ is computed after $c[1, 2]$.





$c[i, j]$ depends on:

- $c[i, j - 1], c[i, j - 2], \dots, c[i, i]$
- $c[i + 1, j], c[i + 2, j], \dots, c[j, j]$

Fixing the Problem

- Solution 1: We can always use memoization (left as an exercise)
- Solution 2: Bottom-up approach in a different order



Solution (a)

Pseudocode

```
1: procedure DPCost( $p, 1, n$ )
2:    $c[1 \dots n+1, 0 \dots n] \leftarrow$  new array
3:   for  $i \leftarrow 1, n+1$  do                                ▷ base cases
4:      $c[i, i-1] \leftarrow 0$ 
5:   for  $i \leftarrow n, 1$  do                                ▷ recursive case
6:     for  $j \leftarrow i, n$  do
7:        $c[i, j] \leftarrow \infty$ 
8:       for  $r \leftarrow i, j$  do
9:          $c[i, j] \leftarrow \min(c[i, j], c[i, r-1] + c[r+1, j])$ 
10:      for  $k \leftarrow i, j$  do
11:         $c[i, j] \leftarrow c[i, j] + p[k]$ 
12:   return  $c[1, n]$ 
```

Solution (b)

Pseudocode

```
1: procedure DPCost( $p, 1, n$ )
2:    $c[1 \dots n + 1, 0 \dots n] \leftarrow$  new array
3:   for  $i \leftarrow 1, n + 1$  do                                ▷ base cases
4:      $c[i, i - 1] \leftarrow 0$ 
5:   for  $s \leftarrow 0, n - 1$  do                                ▷ recursive case
6:     for  $i \leftarrow 1, n - s$  do
7:        $j \leftarrow i + s$ 
8:        $c[i, j] \leftarrow \infty$ 
9:       for  $r \leftarrow i, j$  do
10:          $c[i, j] \leftarrow \min(c[i, j], c[i, r - 1] + c[r + 1, j])$ 
11:       for  $k \leftarrow i, j$  do
12:          $c[i, j] \leftarrow c[i, j] + p[k]$ 
13:   return  $c[1, n]$ 
```

Comparison

- Solution (a) was easier: We only changed one line of code!
- Solution (b) proceeds by *increasing size* s .
- This is a general approach in bottom-up DP: define a size for the subproblems, and solve the smaller problems first.
- Here the size is $s = j - i$, which is the number of nodes in the tree (minus one).

Analysis

- Both algorithms run in $\Theta(n^3)$ time due to the three nested loops.
(Detailed analysis: see exercise set.)
- Lines 10-11 on Slide 18 take $\Theta(n^3)$ time as well.
- We can do the same in $\Theta(n^2)$ time. (See next slide.)
- The overall running time is still $\Theta(n^3)$, but it should be faster by a constant factor in practice.

Optimized Solution

Precomputing all sums $P[i, j] = \sum_{k=i}^j p_k$

```
1: procedure INITP( $p, 1, n, P$ )
2:   for  $i \leftarrow 1, n$  do
3:      $P[i, i] \leftarrow p[i]$ 
4:     for  $j \leftarrow i + 1, n$  do
5:        $P[i, j] \leftarrow P[i, j - 1] + p[j]$ 
```

- This procedure takes $\Theta(n^2)$ time.

Optimized Solution

Pseudocode

```
1: procedure DPCost( $p, 1, n$ )
2:    $c[1 \dots n + 1, 0 \dots n] \leftarrow$  new array
3:    $P[1 \dots n, 1 \dots n] \leftarrow$  new array
4:   INITP( $p, 1, n, P$ )
5:   for  $i \leftarrow 1, n + 1$  do                                ▷ base cases
6:      $c[i, i - 1] \leftarrow 0$ 
7:   for  $i \leftarrow n, 1$  do                                ▷ recursive case
8:     for  $j \leftarrow i, n$  do
9:        $c[i, j] \leftarrow \infty$ 
10:      for  $r \leftarrow i, j$  do
11:         $c[i, j] \leftarrow \min(c[i, j], c[i, r - 1] + c[r + 1, j])$ 
12:         $c[i, j] \leftarrow c[i, j] + P[i, j]$ 
13:   return  $c[1, n]$ 
```

Computing an Optimal BST

Pseudocode

```
1: procedure DPCost( $p, 1, n, R$ )
2:    $c[1 \dots n + 1, 0 \dots n] \leftarrow$  new array
3:    $P[1 \dots n, 1 \dots n] \leftarrow$  new array
4:   INITP( $p, 1, n, P$ )
5:   for  $i \leftarrow 1, n + 1$  do                                ▷ base cases
6:      $c[i, i - 1] \leftarrow 0$ 
7:   for  $i \leftarrow n, 1$  do                                ▷ recursive case
8:     for  $j \leftarrow i, n$  do
9:        $c[i, j] \leftarrow \infty$ 
10:      for  $r \leftarrow i, j$  do
11:        if  $c[i, r - 1] + c[r + 1, j] < c[i, j]$  then
12:           $c[i, j] \leftarrow c[i, r - 1] + c[r + 1, j]$ 
13:           $R[i, j] \leftarrow r$            ▷ root of optimal subtree  $T^*[i, j]$ 
14:       $c[i, j] \leftarrow c[i, j] + P[i, j]$ 
```

Computing an Optimal BST

Pseudocode

```
1: procedure OPTBST( $p, 1, n, R$ )
2:    $R[1 \dots n, 1 \dots n] \leftarrow$  new array
3:   DPCOST( $p, 1, n, R$ )
4:   return RECBST( $1, n, R$ )
```

Pseudocode

```
1: procedure RECBST( $i, j, R$ )
2:   if  $i > j$  then
3:     return empty tree
4:    $r \leftarrow R[i, j]$ 
5:    $T \leftarrow$  NEWNODE ( $k_r$ )
6:    $\text{left}(T) \leftarrow$  RECBST( $i, r - 1, R$ )
7:    $\text{right}(T) \leftarrow$  RECBST( $r + 1, j, R$ )
8:   return  $T$ 
```

Computing an Optimal BST

- OPTBST takes $O(n)$ time
- As usual, it is easy to reconstruct a solution if we are able to compute its value, by a simple modification of the algorithm.

Concluding Remarks

- The brute force approach takes exponential time, and dynamic programming takes cubic time.
- So once again dynamic programming improved the running time from exponential to polynomial.
- A running time of $\Theta(n^3)$ suggests that we can solve the problem up to roughly $n = 1000$. The brute force approach probably fails for n less than 50, as we saw in the case of binomial coefficients.
- Difference with the textbook: The textbook proves a more general result where we assign a probability q_i to each interval (k_{i-1}, k_i) . It makes it more technical, but the main idea is the same.